

TemporalGraph-ToM: Evaluating Theory-of-Mind Reasoning with Temporal Belief Graphs

Dijana Tolic, PhD

April 2026

Extended Abstract

Higher-order theory of mind (ToM), the ability to attribute mental states to others, including their beliefs about others’ beliefs, remains a consistent challenge in large language model reasoning. Existing benchmarks evaluate higher-order belief recursion, such as Hi-ToM [1], interactional information asymmetry in FANToM [2], broad social-cognitive ToM abilities in ToMBench [3], first-person social and counterfactual reasoning in EgoSocialArena [4], and comprehensive cognition-inspired ToM paradigms in CogToM [5].

We introduce TemporalGraph-ToM (TGToM), a benchmark of 1k multi-agent stories that probes higher-order theory of mind with temporal belief graphs. TGToM uses controlled simulations up to seventh-order belief attribution, with information asymmetry from public and private claims, observation asymmetry from interleaved exits, agent reasoning involving deception and cooperation, and higher-order counterfactual scenarios. Evaluation questions span world state, higher-order beliefs, causal beliefs, common knowledge, and counterfactual beliefs.

Probing the limits of higher-order reasoning is relevant to LLM safety, agent safety, human-AI interaction, and reasoning interpretability. In evaluations of frontier models, we find a reality-bias (omniscient reasoning) failure mode: models default to the actual world state even when an agent’s belief diverges from it. TGToM evaluates whether models can correctly attribute belief updates under a belief propagation rule and exposes misattribution when models confuse the true world state with an agent’s propagated belief.

Reasoning scaffold. We introduce Temporal Belief Graph (TBG) and compare it with CoT and System 2 prompting. TBG is formalized as a temporal multilayer graph where nodes represent each agent’s state updated after an event. Beliefs propagate through four directed link types in two layers:

communication links (trusted claim, untrusted claim) and intent links (cooperative, deceptive). Beliefs additionally update directly when an agent observes an event. The node states and link types are generalizable to multiple beliefs, other mental states, event types, and agent interactions.

Temporal Belief Graph

Temporal belief graph is formalized as a temporal multilayer graph where nodes represent each agent’s state (belief, intent, observation) updated after an event. Node state for agent i at time t :

$$s_i^{(t)} = \left(b_i^{(t)}, \iota_i^{(t)}, o_i^{(t)} \right)$$

where:

- $b_i^{(t)}$: belief
- $\iota_i^{(t)}$: intent
- $o_i^{(t)}$: observation

Beliefs propagate through four directed link types in two layers: communication links (trusted claim, untrusted claim) and intent links (cooperative, deceptive). Beliefs additionally update directly when an agent observes an event. Directed link types:

$$e_{j \rightarrow i}^{(t)} \in \{\text{trusted, untrusted, cooperative, deceptive}\}$$

Update rule given event:

$$s_i^{(t+1)} = \begin{cases} f_{\text{obs}}(s_i^{(t)}, E_t) & \text{if observed} \\ f_{\text{trust}}(s_i^{(t)}, E_t) & \text{if trusted} \\ f_{\text{intent}}(s_i^{(t)}, E_t) & \text{if deceptive} \\ s_i^{(t)} & \text{otherwise} \end{cases}$$

Related LLM reasoning scaffolds include temporal belief-state chains, as in TimeToM [6], and graph-structured reasoning, as in Graph of Thoughts (GoT) [7]. Temporal Belief Graph combines the temporal dimension with graph representations for multi-agent belief propagation.

References

- [1] Yinghui He, Yufan Wu, Yilin Jia, Rada Mihalcea, Yulong Chen, and Naihao Deng. Hi-ToM: A Benchmark for Evaluating Higher-Order Theory of Mind Reasoning in Large Language Models. Findings of EMNLP 2023.
- [2] Hyunwoo Kim et al., FANToM: A Benchmark for Stress-testing Machine Theory of Mind in Interactions. EMNLP 2023.
- [3] Zhuang Chen et al., ToMBench: Benchmarking Theory of Mind in Large Language Models. ACL 2024.
- [4] EgoSocialArena, Guangyao Hou et al., EgoSocialArena: Benchmarking the Social Intelligence of Large Language Models from a First-person Perspective. 2024.
- [5] CogToM, Haibo Tong, Zeyang Yue, Feifei Zhao, Erliang Lin, Lu Jia, Ruolin Chen, Yinqian Sun, Qian Zhang, and Yi Zeng. CogToM: A Comprehensive Theory of Mind Benchmark inspired by Human Cognition for Large Language Models. 2026.
- [6] Guiyang Hou, Wenqi Zhang, Yongliang Shen, Linjuan Wu, and Weiming Lu. TimeToM: Temporal Space is the Key to Unlocking the Door of Large Language Models’ Theory-of-Mind. Findings of ACL 2024.
- [7] Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, M. Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Niewiadomski, Piotr Nyczyk, and Torsten Hoefer. Graph of Thoughts: Solving Elaborate Problems with Large Language Models. AAAI 2024.